

Parallel programming

Introduction





Why should you care about it?

- Sometimes you **want to get the result faster** – the algorithm with big amount of computation / big amount of data

Application: scientific world (*simulations, calculations*), big data computing (*faster processing, databases*), machine learning, deep learning



- Sometime you **have limited time to fulfill task**, sequential way is too slow – real time processing
- Benefit: some general principles are applicable in thinking about architecture of separate programs over related tasks




Why should you care about it?

- Parallel computing is a dominant player in scientific and cluster computing. Why?
 - Moore law (number of transistors doubles about every two years; for same price, price per power halving) is reaching its limits
 - Increase in transistor density is limited
 - Memory access time has not been reduced at a rate comparable with processing speed

Moore's Law: The number of transistors on microchips doubles every two years
 Moore's law describes the empirical regularity that the number of transistors on integrated circuits doubles approximately every two years. This advancement is important for other aspects of technological progress in computing – such as processing speed or the price of computers.

Our World
in Data

Transistor count
50,000,000,000




2. This chart shows the Linley Group's "Cost Per Transistor" curve (2017) taken from Cadence's "Breakfast Bytes" blog.



Why should you care about it?


- How to get out of this trap?
 - Most promising approach is to have multiple cores on a single processor
 - Number of cores growing, speed per core growing slower
 - Today's desktop computers 2023 offer
 - Intel Core i9-13900KS - 24 cores, 32 threads, 3,2GHz (TDP 253W), Boost 6 GHz.
 - AMD Ryzen 9 7950X3D - 16 cores, 32 threads, 4,2GHz (TDP 120W), Boost 5,7 GHz
 - Parallel computing can be found at many devices today:





Ok; However, It should be task for compiler and not for me!!!

- Yes, compiler can help you, but without your guidance, it is not able pass all the way to the successful result.
 - Parallel programs often look very different than sequential ones
 - An efficient parallel implementation of a serial program may not be obtained by simply parallelizing each step
 - Rather, the best parallelization may be obtained by stepping back and devising an entirely new algorithm
 - Instruction level paralelization






What is the aim of labs?

- To get the feel for parallel programming
 - 1) Understand what makes the parallelisation **complicated**
 - 2) Which **problems** can occur during the paralellisation
 - 3) What can be a **bottleneck**
 - 4) How to think about **algorithms** from the paralellisation point of view



Familiar terms: race condition, false sharing, synchronizaton, deadlocks, communication overhead, work disbalance, idling, another design of algorithm vs. sequential version

- To get basic skills in common parallel programming frameworks
 - 1) for Multicore processors
 - 2) for Computer clusters
 - 3) for GPU (nice opportunity to play with)





Seminar topics

- OpenMP – for Multicore processors, easy way to parallelize originally sequential code, UMA concept
- MPI – for Computer clusters, concept of units communicating through messages, NUMA concept
- Numba – computation on GPU
- Theoretical seminars – helps to prepare for the exam




Course web

- Course page <https://cw.fel.cvut.cz/b231/courses/pag/start>
 - Detailed plan of the labs, grading



What does this course require?


- Knowledge of C, C++, basics of Python



- Analytical thinking and being open-minded



- Basic skills with Linux – shell, ssh, etc. (for MetaCentrum)





Setting up

- Installation at home
- Be prepared for coding next week
- Small helloworld examples prepared for you to check if enviroment runs smoothly
- Recommendations follow





Our recommendations



Linux, Mac OS, Windows

- CMake and g++
- Recommended IDE: **CLion**
 - <https://download.cvut.cz>, JetBrains
- Homework and semestral project skeletons provided only as Cmake projects
- See next slides for your platform

Windows+Visual Studio? :(

- Use at your own risk
- **Do not use MSVC** (no support for newer OpenMP)



Ubuntu toolchain

- You can use inofficial PPA for the Clion, see [this link](#).

- Install g++ and cmake

```
>> sudo apt install g++ cmake [gdb]
```

- Install MPI library

```
>> sudo apt install libopenmpi-dev
```



Windows mingw toolchain


- Install msys2, see [this link](#)
- In the msys2 console do the following

```
>> pacman -Syu
>> pacman -Su
>> pacman -S base-devel mingw-w64-x86_64-toolchain
>> pacman -S mingw-w64-x86_64-msmpi
```
- Create MinGW toolchain in CLion, see [this link](#). If msys2 is installed in default location, set `C:\msys64\mingw64` as your MinGW Environment path (everything else should be detected automatically), Setup in CLion Settings -> Cmake -> Generator on value MINGW Makefiles
- Add msys2 directories to your PATH environment variable, e.g.,

```
C:\msys64
C:\msys64\mingw64\bin
```
- If MPI library found, but program returns nonzero code and no output printed, try install [this link](#), magically helped



Mingw CLion settings CMake



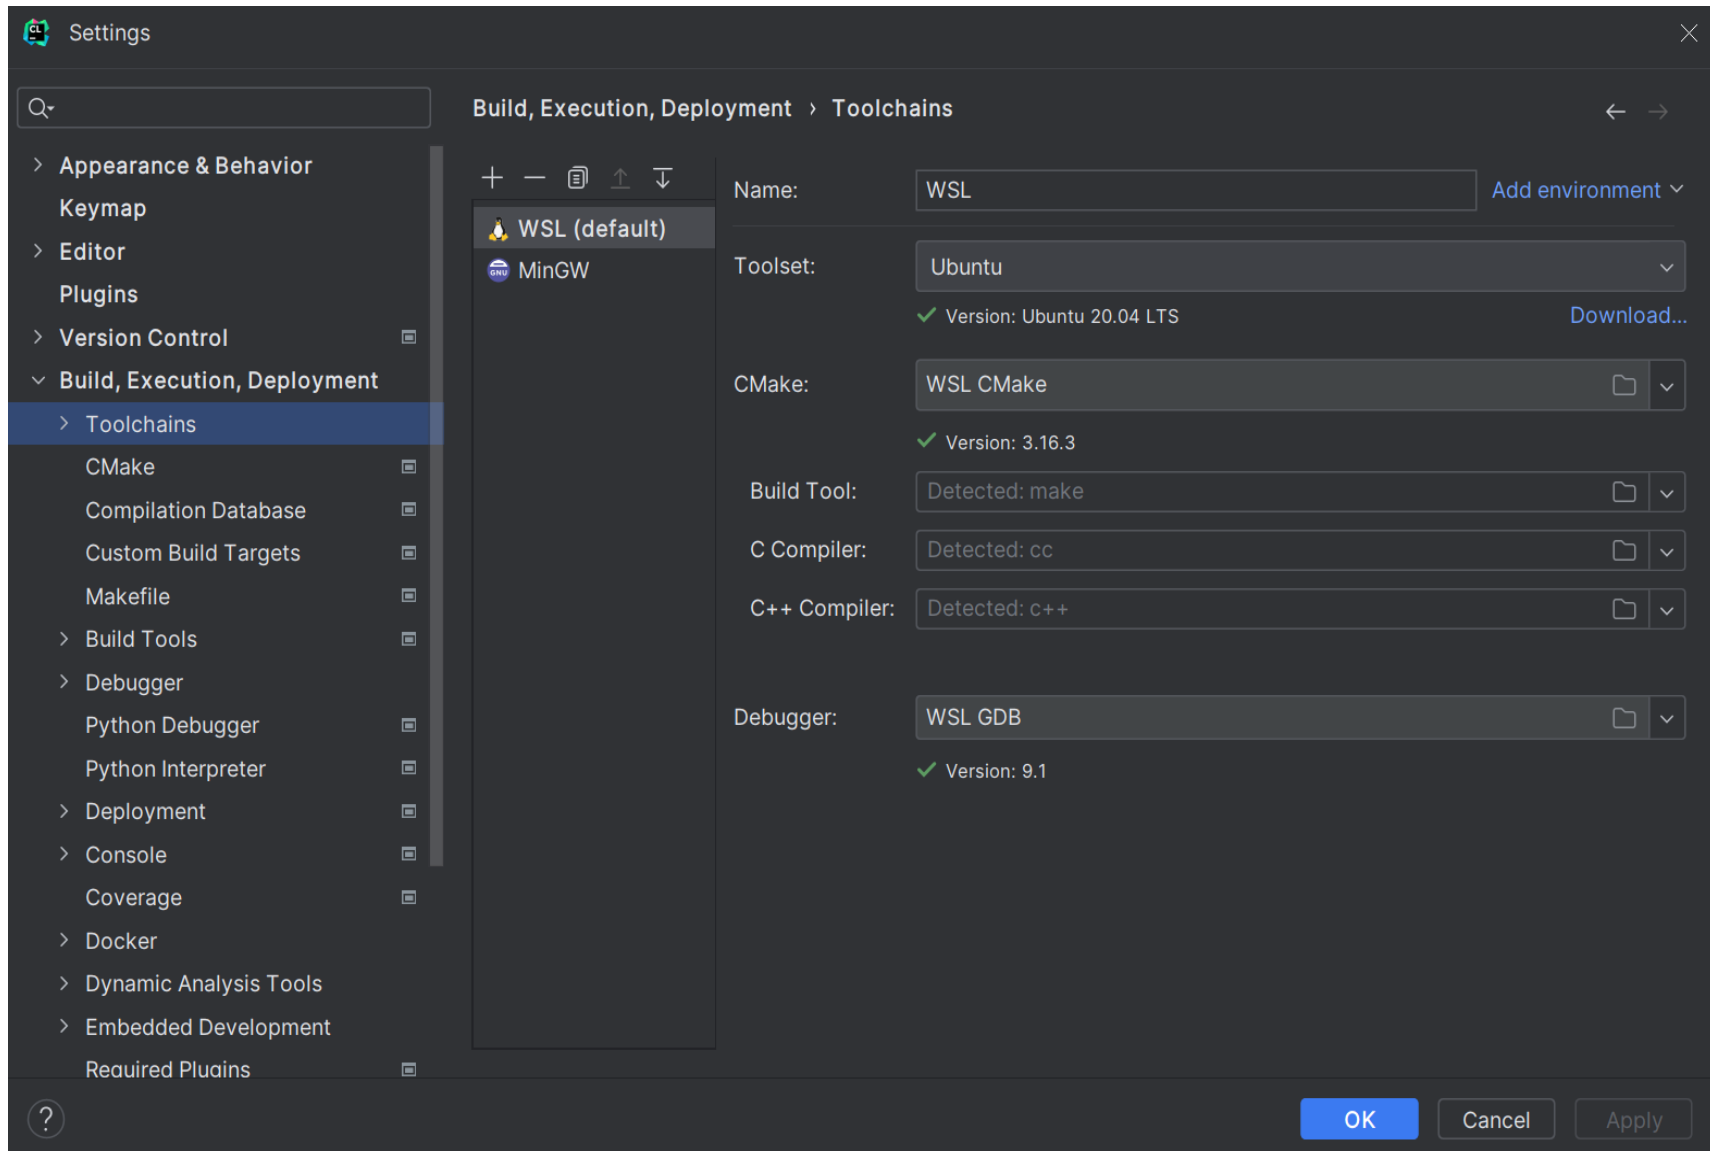


Windows alternative WSL toolchain

- Install WSL, see [this link](#)
 - In powershell run: `wsl --install`
- Install Ubuntu distribution via microsoft store, see [this link](#)
- Open Ubuntu terminal, initiate system (user access setup, first run), install following
 - `sudo apt-get update`
 - `sudo apt install g++ cmake gdb`
 - `sudo apt install libopenmpi-dev`
- Set up WSL in Clion toolchains, see [this link](#)

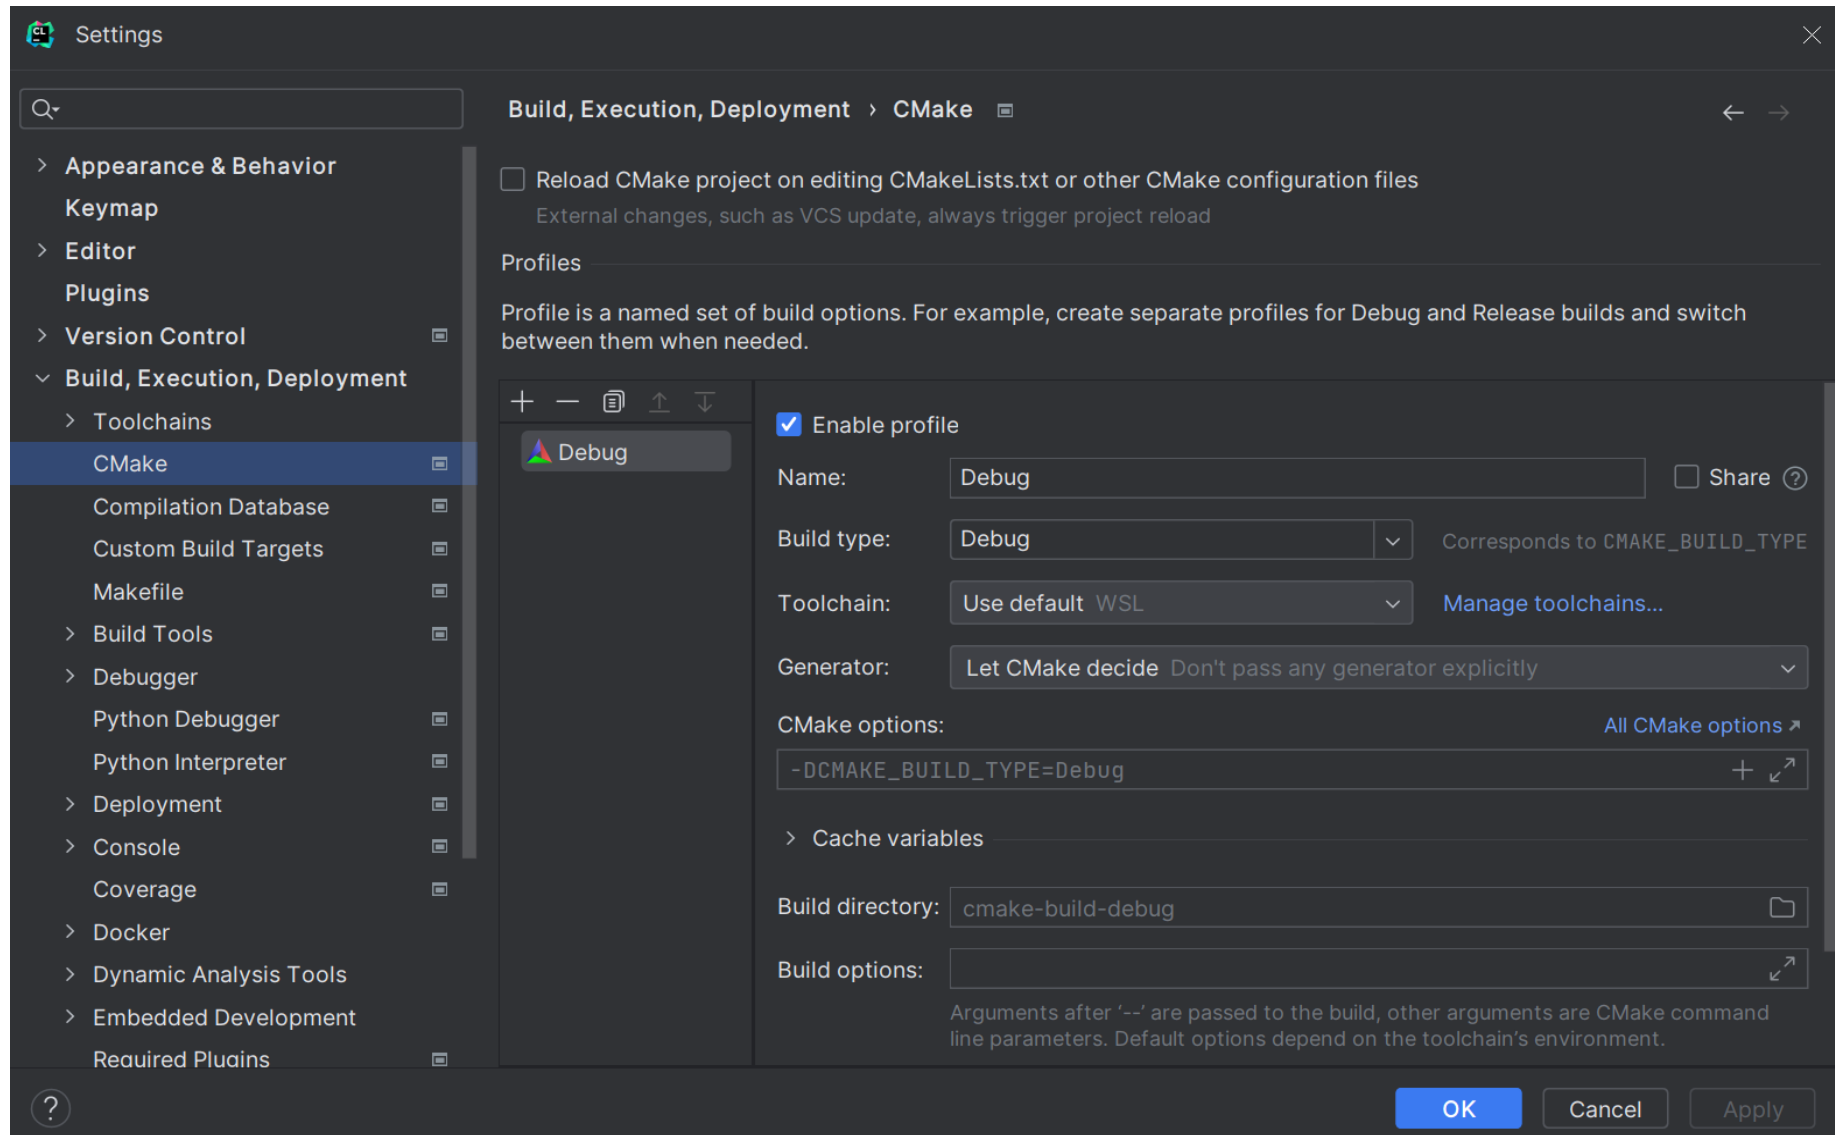


WSL CLion settings Toolchains





WSL CLion settings CMake





MacOS toolchains

- Using g++ (**recommended**)
 - Install g++ from Homebrew
`>> brew install gcc`
 - Find the installed g++ executable. Usually a program called **g++-FÑ** where **FÑ** is the version (can be found using TAB completion), e.g., **g++-9**
 - Set **g++-FÑ** compiler in CLion: Settings → Build, execution, Deployment → Toolchains → C++ compiler
- Using clang
 - Install OpenMP runtime from Homebrew
`>> brew install libomp`
 - Check where libomp is installed, usually `/usr/local/opt/libomp`
`>> brew --prefix libomp`
 - Link OpenMP into CMakeLists.txt
`include_directories("/usr/local/include" "/usr/local/opt/libomp/include")`
`link_directories("/usr/local/lib" "/usr/local/opt/libomp/lib")`
- Install MPI
`>> brew install open-mpi`



Expected cmake console print for all

- Found OpenMP/MPI TRUE
OPENMP (needed for next week), MPI (there is time to solve issues, used later)

```
CMake  Debug
-- Check for working CXX compiler: /usr/bin/c++ -- works
-- Detecting CXX compiler ABI info
-- Detecting CXX compiler ABI info - done
-- Detecting CXX compile features
-- Detecting CXX compile features - done
-- Found OpenMP_C: -fopenmp (found version "4.5")
-- Found OpenMP_CXX: -fopenmp (found version "4.5")
-- Found OpenMP: TRUE (found version "4.5")
-- Found MPI_C: /usr/lib/x86_64-linux-gnu/openmpi/lib/libmpi.so (found version "3.1")
-- Found MPI_CXX: /usr/lib/x86_64-linux-gnu/openmpi/lib/libmpi_cxx.so (found version "3.1")
-- Found MPI: TRUE (found version "3.1")
-- Configuring done
-- Generating done
-- Build files have been written to: /mnt/c/Users/stejs/Desktop/week1_codes/cmake/cmake-build-debug

[Finished]
```



Expected program outputs

- Expected console outputs for provided helloworld programs to test your environment

- **OpenMP**

Value of x: 550

Process finished with exit code 0

- **MPI**

My ranking hello world example: 0

Total number of processes: 1

Process finished with exit code 0